

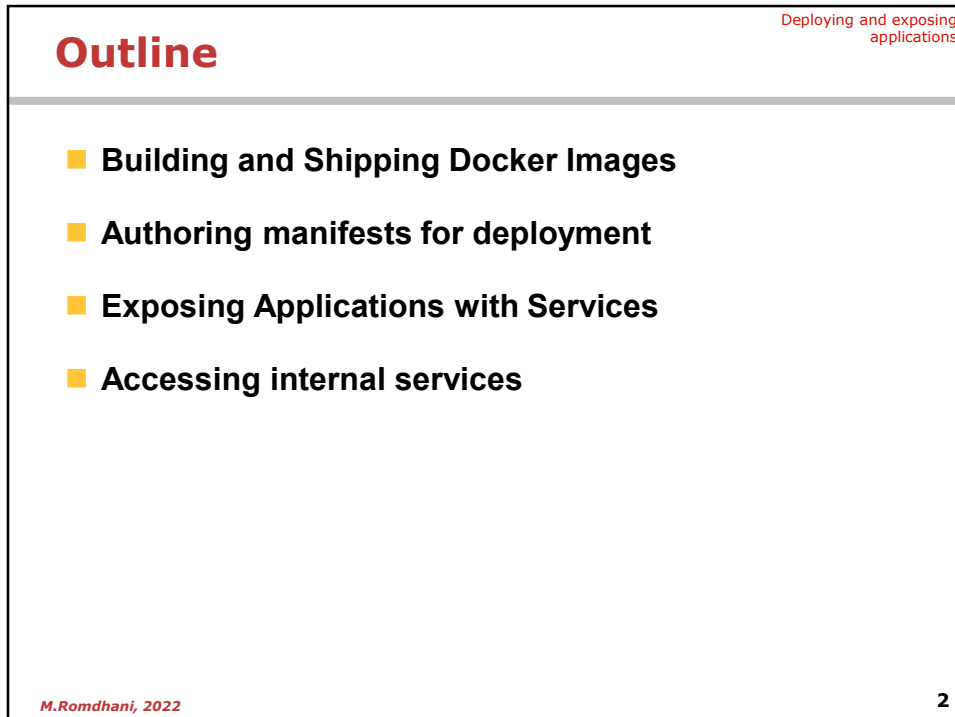
The slide features a purple header and footer bar. The header bar contains a small collage of images on the left. The main content area is white and contains the Docker logo (a blue hexagon with a white ship's wheel) at the top center. Below the logo, the text "Unit 2" is written in red. The main title "Deploying and Exposing Applications" is written in a large, bold, red font. In the bottom right corner, there are three small icons (a circle, a square, and a triangle) above the text "Business Training".

Unit 2

Deploying and Exposing Applications

Business Training

1



The slide has a white background. The title "Outline" is written in a large, bold, red font. To the right of the title, the text "Deploying and exposing applications" is written in a smaller, red font. Below the title, there is a list of four items, each preceded by a yellow square bullet point. The items are: "Building and Shipping Docker Images", "Authoring manifests for deployment", "Exposing Applications with Services", and "Accessing internal services". At the bottom left, the text "M.Romdhani, 2022" is written in a small, red font. At the bottom right, the number "2" is written in a small, red font.

Outline

Deploying and exposing applications

- Building and Shipping Docker Images
- Authoring manifests for deployment
- Exposing Applications with Services
- Accessing internal services

M.Romdhani, 2022

2

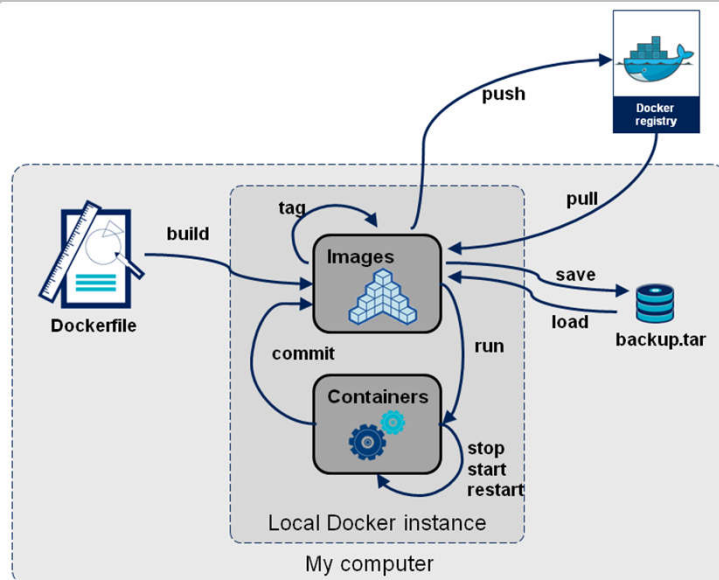
2

Building and Shipping Docker Images

3

Docker containers & images

Deploying and exposing applications



M.Romdhani, 2022

4

4

docker build example

Deploying and exposing
applications

- This **Dockerfile** for creating a simple Java application:

```
FROM maven:3.5.2-jdk-9
COPY src /usr/src/app/src
COPY pom.xml /usr/src/app
RUN mvn -f /usr/src/app/pom.xml clean package

EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/usr/src/app/target/myapp-1.0.0-SNAPSHOT.jar"]
```

- The image can be built with the following command:

```
$ docker build -t my-app .
```

- Notes:

- The build has the current directory as context
- Each command in the Dockerfile creates a new (temporary container)
- Every creation step generates a layer that is cached, so repeated builds are fast

M.Romdhani, 2022

5

5

Dockerfile instructions

Deploying and exposing
applications

Instruction	Description
FROM	Parent image
ARG	Parameters for constructing the image
ENV	Specify Environment variables
LABEL	Specify Label meta-data
VOLUME	Mount volumes
RUN	Run a command
COPY	Copy files to the image
ADD	Add files to the image
WORKDIR	Specify the working directory
EXPOSE	Expose ports to be accessed
USER	User name or UID to be used
ONBUILD	Instructions to execute when constructing child images
CMD	Command to execute when starting a container
ENTRYPOINT	The default entry point of the container

M.Romdhani, 2022

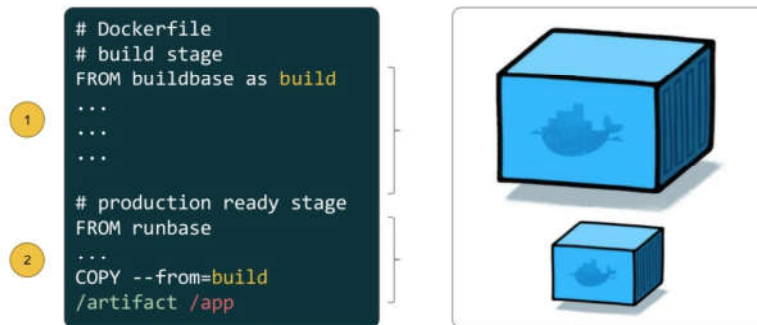
6

6

What are multi stage builds ?

Deploying and exposing applications

- **Multi-stage builds are a method of organizing a Dockerfile to minimize the size of the final container.**
 - This is made possible by the image building process into multiple stages
 - Each stage is a separate image, and can copy files from previous stages.



M.Romdhani, 2022

7

7

Multi-stage builds in practice

Deploying and exposing applications

- **Building a Java Spring Boot Application in a single image**

```

FROM maven:3.5.2-jdk-9
COPY src /usr/src/app/src
COPY pom.xml /usr/src/app
RUN mvn -f /usr/src/app/pom.xml clean package

EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/usr/src/app/target/myapp-1.0.0-SNAPSHOT.jar"]
  
```

- **Building the Java Spring Boot Application using a multi-stage build**

```

FROM maven:3.5.2-jdk-9 AS build
COPY src /usr/src/app/src
COPY pom.xml /usr/src/app
RUN mvn -f /usr/src/app/pom.xml clean package

FROM openjdk:9-jre-alpine
COPY --from=build /usr/src/app/target/ myapp-1.0.0-SNAPSHOT.jar
/usr/app/myapp-1.0.0-SNAPSHOT.jar
EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/usr/app/myapp-1.0.0-SNAPSHOT.jar"]
  
```

M.Romdhani, 2022

8

8

Useful Docker Commands

Deploying and exposing applications

- **docker ps** list running containers.
- **docker ps -a** list all container including stopped container
- **docker pull** download a image from Docker Hub registry.
- **docker build** . to build a container based on the Dockerfile in the current directory (the dot). **docker build -t myimage:latest** . creates a container and stores the image under the given name
- **docker images** or **docker image ls** shows all local storage images
- **docker run** to run a container using the image given in parameter
- **docker logs** display the logs of a container, you specified. To continue showing log updates just use **docker logs -f mycontainer**
- **docker volume ls** lists the volumes, which are commonly used for persisting data of Docker containers.
- **docker network ls** - list all networks available for docker container
- **docker network connect** adds the container to the given container network. That enables container communication by simple container name instead of IP.
- **docker rm** removes one or more containers. **docker rm mycontainer**, but make sure the container is not running
- **docker rmi** removes one or more images. **docker rmi myimage**, but make sure no running container is based on that image
- **docker stop** stops one or more containers.

M.Romdhani, 2022

9

9

Shipping an Image to a registry

Deploying and exposing applications

- **An account is required to push images to Dockerhub**
- **Ship an image to DockerHub.com**
 - > **docker push myrepo/myimage:1.0**
- **DockerHub has a limit of the number of Pulls, since 11/2020**
 - 100 pulls in 6 Hours /Anonymous
 - 200 pulls in 6 Hours/Logged in user
- **Alternatives to DockerHub**
 - RedHat Quay.io
 - Amazon Elastic Container Registry (ECR)
 - JFrog Artifactory.
 - Azure Container Registry.
 - Google Container Registry.
 - VMWare Harbor

M.Romdhani, 2022

10

10

Dockerfile best practices

Deploying and exposing
applications

- Use official base images
- Prefer COPY over ADD
- Group RUN instructions in one line
- ADD a .dockerignore file
- Use Multi-stage builds

M.Romdhani, 2022

11

11

Authoring manifests for deployment

12

Yaml manifest structure

Deploying and exposing applications

Required fields

- **apiVersion** - Which version of the Kubernetes API you're using to create this object
- **kind** - What kind of object you want to create
- **metadata** - Data that helps uniquely identify the object, including a name string, UID, and optional namespace
- **spec** - What state you desire for the object. The precise format of the object spec is different for every Kubernetes object

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

The status field

- While spec describes the desired state, the status describes the current state. It is added and updated continuously by K8s control plane.
- kubectl get deploy mydepl -o yaml

M.Romdhani, 2022

13

13

kubectl apply vs create

Deploying and exposing applications

- **kubectl create -f whatever.yaml**
 - creates resources if they don't exist
 - if resources already exist, don't alter them (and display error message)
- **kubectl apply -f whatever.yaml**
 - creates resources if they don't exist
 - if resources already exist, update them (to match the definition provided by the YAML file)
 - stores the manifest as an annotation in the resource

M.Romdhani, 2022

14

14

Simple Pod Deployment

Deploying and exposing applications

Deployment steps

1. Describe the app using Kubernetes YAML (**my-nginx-pod.yaml**)
2. Run the deployment Command
kubectl apply -f my-nginx-pod.yaml
3. Make sure the pod has been created
kubectl get pods
4. Tear down your app
kubectl apply -f my-nginx-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: mynginxapp
  labels:
    name: mynginxapp
spec:
  containers:
    - name: mynginxapp
      image: nginx
      ports:
        - containerPort: 80
```

M.Romdhani, 2022

15

15

Simple Pod with namespace and labels

Deploying and exposing applications

Additional information

- **Namespace:** Namespaces provide a scope for Kubernetes resources, splitting the cluster in smaller units.
- **Labels:** Labels are intended to be used to specify identifying attributes of objects that are meaningful and relevant to users, but do not directly imply semantics to the core system.

```
apiVersion: v1
kind: Pod
metadata:
  name: mynginxapp
  namespace: default
  labels:
    name: mynginxapp
    profile: dev
spec:
  containers:
    - name: mynginxapp
      image: nginx
      ports:
        - containerPort: 80
```

M.Romdhani, 2022

16

16

A Multi container Pod: Main Container with Side Car Container

Deploying and exposing applications

```

apiVersion: v1
kind: Pod
metadata:
  name: pod-with-sidecar
spec:
  # Create a volume called 'shared-logs' that the pp and sidecar share.
  volumes:
    - name: shared-logs
      emptyDir: {}
  containers:
    - name: app-container # Main application container
      # Simple application: write the current date to the log file every 5 seconds
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do date >> /var/log/app.txt; sleep 5;done"]
      volumeMounts: # Mount the pod's shared log file into the app container
        - name: shared-logs
          mountPath: /var/log

    - name: sidecar-container # Sidecar container
      image: nginx:1.7.9
      ports:
        - containerPort: 80
      volumeMounts: # Mount the pod's shared log file into the sidecar
        - name: shared-logs
          mountPath: /usr/share/nginx/html # nginx-specific mount path

```

■ Main Container and the Side Car Container share a Volume

M.Romdhani, 2022

17

17

Using Deployments

Deploying and exposing applications

■ Saving this manifest into nginxdeploy.yaml and submitting it to a Kubernetes cluster will create the defined Deployment, ReplicaSet and the Pods

- You can then get the current Deployments deployed:
kubect1 get deployments
- You can then get the current ReplicaSets deployed:
kubect1 get rs
- You can then get the current pods deployed:
kubect1 get pods

```

# for versions before 1.9.0 use apps/v1beta2
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80

```

M.Romdhani, 2022

18

18

Updating the deployment

Deploying and exposing applications

- You can update the deployment by applying a new YAML file. This YAML file specifies that the deployment should be updated to use nginx 1.16.1

```
# for versions before 1.9.0 use apps/v1beta2
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods
  template: Pod Template
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.16.1
          ports:
            - containerPort: 80
```

M.Romdhani, 2022

19

19

Scaling the application by increasing the replica count

Deploying and exposing applications

- You can increase the number of pods in your Deployment by applying a new YAML file. This YAML file sets replicas to 4, which specifies that the Deployment should have four pods:

```
# for versions before 1.9.0 use apps/v1beta2
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 4 # Update the replicas from 2 to 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.16.1
          ports:
            - containerPort: 80
```

M.Romdhani, 2022

20

20

Exposing Applications with Services

21

Services

Deploying and exposing applications

- Services give us a **stable endpoint** to connect to a pod or a group of pods
 - Durable resource (unlike Pods)
 - static cluster-unique IP
 - Target Pods using equality **based selectors**
 - kube-proxy provides simple load-balancing.
- A Kubernetes Service can select the pods it is supposed to abstract through a **label selector**
- We can create a service either using the command **kubectl expose** or using a Yaml manifest
 - Services are automatically added to an internal DNS zone
- A service has a number of "endpoints"

M.Romdhani, 2022

22

22

Service Types

Deploying and exposing applications

- There are 3 major service types:
 1. **ClusterIP** (default)
 2. **NodePort**
 3. **LoadBalancer**
- There is also another type of services: **ExternalName**. They do not have selectors and uses DNS names instead. They serve as a way to return an alias to an external service residing outside the cluster.

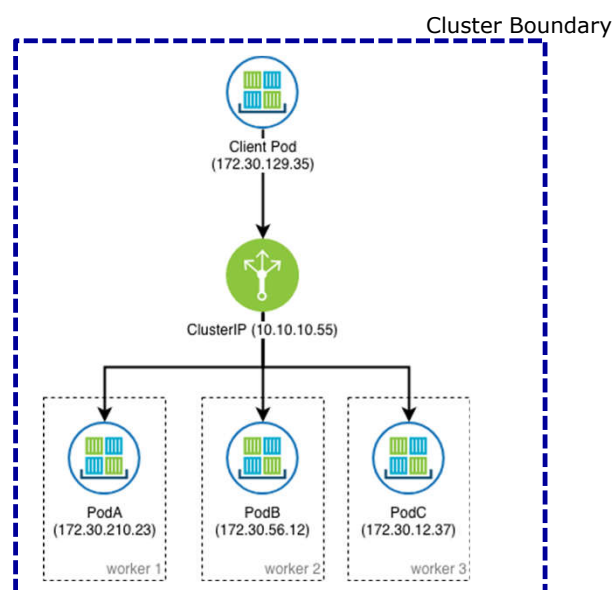
M.Romdhani, 2022

23

23

ClusterIP Services

Deploying and exposing applications



M.Romdhani, 2022

24

24

Deploying and exposing applications

ClusterIP Services

- It is the default service type
- A virtual IP address is allocated for the service
- This IP address is reachable only from **within the cluster (nodes and pods)**
- Perfect for internal communication, within the cluster

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ClusterIP
  selector:
    app: nginx
    env: prod
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80 # should meet the Container port
```

25

M.Romdhani, 2022

25

Deploying and exposing applications

NodePort Services

- NodePort services extend the ClusterIP service.
 - Exposes a port on every node's IP.
- Port can either be statically defined, or dynamically taken from a range between 30000-32767.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: NodePort
  selector:
    app: nginx
    env: prod
  ports:
    - nodePort: 30008
      port: 80
      targetPort: 80
```

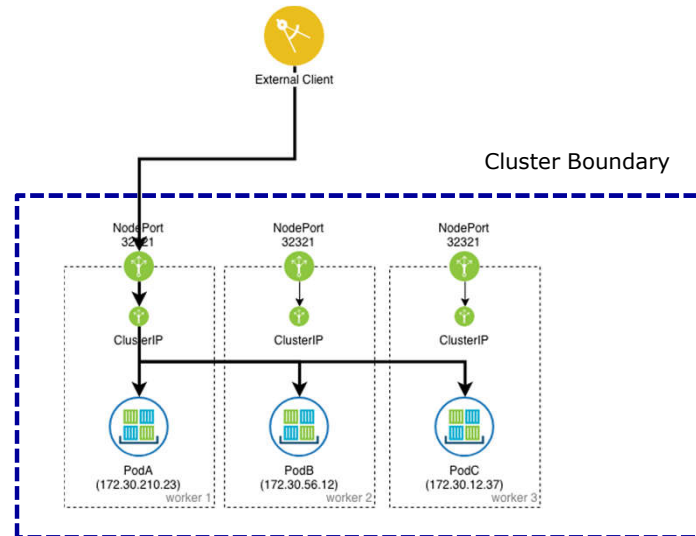
26

M.Romdhani, 2022

26

NodePort Services

Deploying and exposing applications



M.Romdhani, 2022

27

27

nodePort, port, targetPort

Deploying and exposing applications



```
service-definition.yml
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
```

M.Romdhani, 2022

28

28

LoadBalancer Services

Deploying and exposing applications

- LoadBalancer services extend NodePort.
- Works in conjunction with an external system to map a cluster external IP to the exposed service (typically a cloud load balancer, e.g. ELB on AWS, GLB on GCE ...)

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    env: prod
  ports:
    protocol: TCP
    port: 80
    targetPort: 80
```

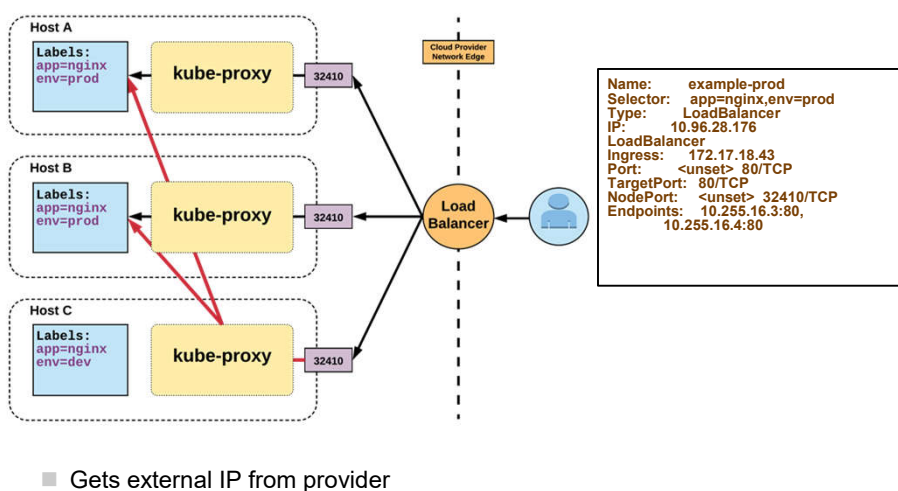
M.Romdhani, 2022

29

29

LoadBalancer Services

Deploying and exposing applications



M.Romdhani, 2022

30

30

Accessing Internal Services

31

Accessing internal services

Deploying and exposing applications

- **When we are logged in on a cluster node, we can access internal services**
 - As per the Kubernetes network model: all nodes can reach all pods and services)
- **When we are accessing a remote cluster, our local machine won't have access to the cluster's internal subnet. To overcome this:**
 - **kubect proxy**: gives us access to the API, which includes a proxy for HTTP resources
 - **kubect port-forward**: allows forwarding of TCP ports to arbitrary pods, services, ...

M.Romdhani, 2022

32

32

kubectl proxy

Deploying and exposing applications

- Running **kubectl proxy** gives us access to the entire Kubernetes API
 - The API includes routes to proxy HTTP traffic
 - By default, the proxy listens on port 8001
- These routes look like the following:
 - /api/v1/namespaces/<namespace>/services/<service>/proxy
- We just add the URI to the end of the request, for instance:
 - /api/v1/namespaces/<namespace>/services/<service>/proxy/index.html
- We can access services and pods this way !
- **Security considerations : kubectl proxy is intended for local use**
 - Running kubectl proxy openly is a huge security risk
 - It is slightly better to run the proxy where you need it (and copy credentials, e.g. ~/.kube/config, to that place)
 - It is even better to use a limited account with reduced permissions

M.Romdhani, 2022

33

33

kubectl port-forward

Deploying and exposing applications

- What if we want to access a TCP service?
 - We can use **kubectl port-forward** instead
 - It will create a TCP relay to forward connections to a specific port (of a pod, service, deployment...)
- The syntax is:


```
kubectl port-forward service/name_of_service local_port:remote_port
```

 - If only one port number is specified, it is used for both local and remote ports

M.Romdhani, 2022

34

34

Additional Deployment Objects

35

DaemonSets

Deploying and exposing applications

■ **A DaemonSet ensures that all (or some) Nodes run a copy of a Pod.**

- As nodes are added to the cluster, Pods are added to them.
- As nodes are removed from the cluster, those Pods are garbage collected. Deleting a DaemonSet will clean up the Pods it created.
- Typical uses of a DaemonSet are:
 - Running a cluster storage daemon, such as glusterd, ceph, on each node.
 - Running a logs collection daemon on every node, such as fluentd or filebeat.
 - Running a node monitoring daemon on every node, such as Prometheus Node Exporter

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: my-daemonset
  namespace: my-namespace
  labels:
    key: value
spec:
  template:
    metadata:
      labels:
        name: my-daemonsetcontainer
    ...
  selector:
    matchLabels:
      name: my-daemonsetcontainer
```

M.Romdhani, 2022

36

36

StatefulSets

Deploying and exposing
applications

■ **StatefulSet is the workload API object used to manage stateful applications.**

- StatefulSets represent a set of Pods with unique, persistent identities and stable hostnames that are maintained regardless of where they are scheduled.
- StatefulSets use an ordinal index for the identity and ordering of their Pods.
 - By default, StatefulSet Pods are deployed in sequential order and are terminated in reverse ordinal order.
 - For example, a StatefulSet named web has its Pods named web-0, web-1, and web-2. When the web Pod specification is changed, its Pods are gracefully stopped and recreated in an ordered way; in this example, web-2 is terminated first, then web-1, and so on.
- StatefulSets are suitable for deploying Kafka, MySQL, Redis, ZooKeeper, and other applications needing unique, persistent identities and stable hostnames.

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: k8s.gcr.io/nginx
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
```

M.Romdhani, 2022

37